

A Techno-economic Analysis of Network Migration to Software-Defined Networking

Tamal Das, Marcel Caria and Admela Jukan
Technische Universität Carolo-Wilhelmina zu Braunschweig
Email: {das, caria, jukan} @ida.ing.tu-bs.de

Marco Hoffmann
Nokia Siemens Networks
Email: marco.hoffmann@nsn.com

Abstract—As the Software-Defined Networking (SDN) paradigm gains momentum, every network operator faces the obvious dilemma – when and how to migrate from existing IP routers to SDN-compliant equipments. A single-step complete overhaul of a fully functional network is impractical, while at the same time, the immediate benefits of SDN are obvious. A viable solution is thus a *gradual* migration over time, where questions of which routers should migrate first, and whether the order of migration makes a difference, can be analyzed from techno-economic and traffic engineering perspective. In this paper, we address these questions from the techno-economic perspective, and establish the importance of *migration scheduling*. We propose optimization techniques and greedy algorithms to plan an effective migration schedule, based on various techno-economic aspects, such as technological gains in combinations with CapEx limitations. We demonstrate the importance of an effective migration sequence through two relevant network management metrics, namely, number of alternative paths availed by a node on migration, and network capacity savings. Our results suggest that the sequence of migration plays a vital role, especially in the early stages of network migration to SDN.

I. INTRODUCTION

Software Defined Networking (SDN) is an emerging networking architecture paradigm that separates the control plane from the data plane, facilitating programmability of the network control functions [1]. SDN is expected to reduce the network OpEx by simplifying operations, optimizing resource usage through centralized data/algorithms, and simplifying network software upgrades. It also significantly cuts down a network operator's CapEx, since a commercial-off-the-shelf (COTS) server with a high-end CPU is much cheaper than a high-end router [2]. As the SDN paradigm gains momentum, the migration from existing IP routers to SDN-compliant equipment, e.g., OpenFlow switches, is becoming eminent. In data centers, SDN can be already fully integrated into the network architecture, by migrating the switching and routing infrastructure entirely to SDN. For a medium-to large-scale ISP, on the other hand, a viable approach is to gradually migrate to SDN, for instance, over a multi-period cycle spanning couple of years.

In cases where due to operational and economic considerations not all routers in a fully functional network can be replaced at once, the operator would need to know which network nodes should be migrated first, and which ones later. Some operators may choose to migrate a fixed number of routers in their network, say quarterly, until complete migration of the network. Others may have a hard limit on their CapEx investment available for migration every quarter, and may choose to quarterly migrate variable number of routers within their CapEx limit. In either case, the operators need to understand how their network can make best use of traffic engineering capabilities enabled by SDN, during various stages of a migration process. Especially those stages are of interest, wherein the native IP routing, such as OSPF, co-exists with SDN-enabled traffic-engineered routing.

In this paper, we propose a techno-economic analysis of network migration to software-defined networking, which includes aspects of migration scheduling. To this end, we design the corresponding optimization model, and propose effective greedy algorithms for the same. Although a scheduled migration must intuitively be better than a random one, as our study confirms, our approach exhibits low computational complexity and high degree of optimality, which makes it highly practical. We consider two scenarios of migration, one where an Internet Service Provider (ISP) has a limit on the number of routers migrated per time-step, and the other where it has a limit on the amount of CapEx investment required for migration per time-step. Thus, our main contributions in this paper are two-fold - (a) simultaneous consideration of technological gains and economic viability of migration to SDN, and (b) design of novel greedy algorithms that come remarkably close to the optimal, measured in terms of traffic engineered metrics, such as number of alternative routes, as well as network capacity savings.

The rest of this paper is organized as follows. Section II discusses the related work. Section III presents the reference migration scenario. In Section IV, we present our migration models, discuss their complexity and optimality, and evaluate the same using simulations in Section V. Section VI concludes our paper.




Fig. 1. Reference migration scenario

II. RELATED WORK

Network migration has been studied using *system-dynamics* and *agent-based* models. In the system dynamics approach, the migration problem is treated as a dynamic system [3], where the rate of migration depends on number of migrated agents in the system. On the other hand, in an agent-based approach [4], the system consists of an ensemble of agents, each trying to increase its own utility. Such studies have mostly been conducted for IPv6 [5] and secure BGP [6].

A choice of the prediction interval for the demand and cost forecast has been studied in [7], where it was shown that short prediction intervals may not be able to sufficiently account for future evolutions, while long prediction intervals may be uncertain. An efficient ant colony meta heuristic for the multi-layer, multi-period migration problem formulated as a path-finding problem is proposed in [8]. Stochastic programming approaches have also been proposed to better handle the uncertainties, as in [9]. Timing issues of migration to a new technology have also been studied, and it was shown that demand growth, migration cost, and cost savings from the new technology must be considered [10]. In [11], the number and location of SDN controllers is studied, which is an important related aspect, but outside the scope of this paper.

This paper extends our recent studies on network migration. In [12], we proposed an agent-based model to study benefits of joint migration to multiple technologies, on a case study of Path Computation Element (PCE) and SDN. In [13], we defined the SDN *migration scheduling problem*, and found optimal solution. In this paper, however, our focus is entirely on devising computationally inexpensive but effective heuristics to schedule migration to a single technology, with consideration of techno-economic factors, such as limitations on CapEx investment. To this end, we propose novel greedy algorithms and establish their practical impact in terms of computational complexity, run times, network capacity gains, etc.

III. REFERENCE MIGRATION SCENARIO

The reference migration scenario is illustrated in Figure 1, wherein an operator manages a network of 4 nodes each of legacy routers and SDN-compliant routers using a network management system (NMS). The network is in an intermediate stage of migration, thereby comprising of both legacy and SDN routers. Whereas legacy routers run OSPF, the SDN routers are collectively controlled by a centralized OpenFlow controller, which is in turn managed by the NMS in providing the TE capabilities in the sense of network optimization. Here, what SDN does provide is a thorough configurability of the routing of all flows in the network, and thus a much greater solution space for network optimization, which in turn leads to an overall better performance of TE. Furthermore, we assume that the SDN routers are also OSPF-compliant, in that they are capable of exchanging OSPF packets with legacy routers. This feature is fundamental to gradual migration and can either be implemented in the SDN routers, or the SDN controller [14].

To facilitate traffic engineering actions through SDN, we assume the SDN controller to have access to all necessary information including network topology, traffic monitoring, and routing. Without loss of generality, we assume that traffic engineering is performed only once at the beginning of every migration period. Therefore, a reasonable capacity headroom (e.g., maximum link utilization set to 70%) assures that despite the expected traffic growth, all links are loaded fairly below the threshold until the end of the current migration period. Before the migration starts, routing is based on the OSPF protocol only, i.e., traffic is always forwarded on the shortest path based on the destination IP address. Finally, traffic engineering through OSPF link cost (weight) changes is not considered in this paper, which we justify with the known fact that network operators do not commonly deploy OSPF link weights for the reasons of routing stability [15].

It is important to note that our focus in this paper is on IP layer, and not on network technologies with built-in traffic engineering capabilities, e.g., MPLS-TE. The reason behind that is, as we believe, SDN routers will substitute legacy IP routers, whereas other technologies work in the layers below, and do not allow gradual migration; for instance, one cannot deploy a few Carrier Ethernet switches in combination with legacy SONET equipment.

IV. OPTIMIZATIONS AND GREEDY ALGORITHMS

In this section, we present different optimization techniques and greedy algorithms for SDN migration. We consider two scenarios, one where the ISP has a limit on the number of routers migrated per time-step, and the other where it has a limit on the amount of CapEx investment required for migration

per time-step. We first summarize the overall migration scheduling strategy from our previous work [13], and then discuss our novel contributions here.

To migrate a network, for every node pair, we first compute the least-cost path, as well as other paths between the same node pair with hop-length equal to its least-cost path. These are the paths between a node pair that become available on migration to SDN, thereby enabling traffic engineering in the network. We then identify the *key-nodes* on each path, which are those nodes that *must* be SDN-compliant for a path to be used for traffic engineering. Thus, for the least cost path, there are no key nodes, and migration of all key-nodes ensures availability of all possible paths in the network for traffic engineering. The concept of key-node is illustrated using Figure 2, which shows the least-cost path between s and d , as well as two more paths of same hop-length as that of the least-cost path. The link weights are indicated above each link. As defined, there are no key-nodes for the least-cost path $s - a - b - d$. For the path $s - c - e - d$ to be available, s must be SDN-compliant, and thus s is its key-node. Finally, for the path $s - c - h - d$ to be available, both s and c must be SDN-compliant, and thus those are its key-nodes.

The generic approach to compute these key-nodes for a path is summarized in Algorithm 1, which essentially involves computation of the *fork* node between two paths between the same node pair. A path is thus available only after *all* its key-nodes migrate. The subsequent step is thus to choose the sequence of key-nodes for migration, which we attempt using optimization and greedy approaches.

A. Optimization Approaches

We now present Integer Linear Programs (ILP) to determine the optimal sequence of migration of nodes in a network, with and without cost constraints. We base our ILP on the ILP formulation from [13], both for comparison and completeness, and further extend it to include CapEx constraints. The parameters and variables used in our model are summarized in Table I.

Algorithm 1 Key-node computation for a path

Input: Path p
Output: $K_p \leftarrow$ Set of key-nodes of path p
 $K_p \leftarrow \phi$
 $A \leftarrow$ Set of paths of cost lesser than that of a , between source and destination of p
for all $a \in A$ **do**
 $K_p \leftarrow K_p \cup \{\text{last common node between } a \text{ and } p \text{ from source to destination of } p\}$
end for
return K_p




Fig. 2. Illustration of key-nodes of a path

TABLE I
LIST OF NOTATIONS USED

Parameter	Meaning
$0 \leq t \leq T$	Migration time-step t , number of time-steps T
$p \in P$	Path p , set of all pre-calculated paths P
$n \in N$	Network node n , set of all nodes N
φ_p	Priority of path p
c_n	Cost of migrating node n
C	Total CapEx required to migrate all nodes
α_p	Number of key nodes on path p
β_p^n	Boolean routing parameter, true if node n is a key node on path p
Variable	Meaning
μ_t^n	Boolean to determine if node n is migrated till time-step t
π_t^p	Boolean to determine if path p is available in time-step t

Here, we attempt to maximize the cumulative number of paths available for traffic engineering during the course of migration as our objective, i.e.,

$$\text{Maximize } \sum_t \sum_p \pi_t^p \cdot \varphi_p$$

The availability of path p in time-step t depends on whether all of its key-nodes (α_p in number) have migrated till that time-step. Thus,

$$\forall p, t: \quad \pi_t^p \cdot \alpha_p \leq \sum_n \beta_p^n \cdot \mu_t^n$$

In the case without CapEx restrictions, the total number of nodes that may migrate in a time-step is bounded by the total number of nodes in the network averaged over the entire migration period. Whereas, in the case with CapEx restriction on each time-step, the average CapEx per time-step invested till time-step t must not exceed the average CapEx per time-step over the entire migration period. Thus,

$$\forall t: \sum_n \mu_t^n \leq t \left\lceil \frac{|N|}{T} \right\rceil \quad \text{or} \quad \forall t: \sum_n \mu_t^n \cdot c_n \leq t \left(\frac{C}{T} \right)$$

(without cost) (with cost)

Finally, we restrict a node that has once migrated from *reverse-migrating*, i.e.,

$$\forall t: \quad \mu_t^n \leq \mu_{t+1}^n$$

B. Greedy Algorithms

In our greedy approach to determine a suitable migration schedule of nodes in a network, we essentially focus on the number of additional alternative paths availed by a node on migration. This is summarized in Algorithms 2, 3 and 4. A preliminary operation to our greedy approach is a one-time computation of least-cost path for each node pair, as well as the paths between each node pair equal in hop-length to the corresponding least-cost path.

At each time-step, we compute the list of all unmigrated key-nodes in the network, and then compute the number of additional alternative paths that can be availed by migration of each of these key-nodes. For example, in Figure 2, if node c migrates, then the path $c-h-e$ will be an additional alternative path available to c for traffic engineering. In the case without CapEx restrictions, at each time-step, we sort the nodes in descending order of the number of alternative paths made available by their migration, and select the first $\lceil N/T \rceil$ nodes for migration. In the case with CapEx limitations on each time-step, we first compute the total CapEx available at current time-step, which is sum of the average CapEx per time-step, and the residual CapEx from all previous time-steps. The residual CapEx results from partial utilization of the CapEx available for a time-step. We then sort the unmigrated key-nodes in decreasing order of the ratio of the number of alternative paths availed by migration of a node and the corresponding cost of migration, and select as many nodes for migration as possible while traversing this sorted list as is allowed by the available CapEx for the current time-step. Algorithm 4 thus selects nodes for migration based on the number of additional paths per unit cost of migration.

For example, say for a network, $m = 2$, and the cost of migration of nodes key-nodes a , b , and c are 5, 8 and 10 units, respectively. If at a particular time-step, a , b , and c can respectively avail of 2, 4 and 3 additional alternative paths on migration, and the CapEx available for this time-step is 15 units, then as per Algorithm 3, b and c are selected for migration, whereas, as per Algorithm 4, a and b should migrate.

C. Computational Complexities

To establish the novelty of our greedy approach over the ILP, we compare the time complexities of both approaches. The ILP efficiently scans the entire search space for optimal solutions.

The ILP primarily comprises of a combinatorial problem of partitioning a set S of N elements into its subsets $\{S_i : 1 \leq i \leq T\}$, i.e.,

$$|S| = \prod_{i=1}^{T-1} \binom{z_i}{|S_i|}, \text{ where,}$$

Algorithm 2 Greedy algorithm for migration schedule

Input: $T \leftarrow$ Number of time-steps
for $t \leftarrow 1$ to T **do**
 $U \leftarrow$ List of unmigrated key-nodes in the network

for all $u \in U$ **do**
 $P(u) \leftarrow$ Number of additional alternative paths available by migration of node u
end for
Based on P , use **Algorithm 2** or **3** to select the node(s) to migrate
end for

Algorithm 3 Node selection algorithm without costs

Input: P
 $m \leftarrow$ Number of nodes to migrate per time-step

Sort P in descending order
for $i \leftarrow 1$ to m **do**
Migrate node i
end for

$$z_i = \begin{cases} N & i = 1 \\ N - \sum_{j=1}^{i-1} |S_j| & 1 < i < T \end{cases}$$

We consider all except the last subset to be of equal size $\lceil n/T \rceil$, thereby resulting in the last subset of size $N - (T-1)\lceil N/T \rceil$. Thus,

$$|S_i| = \begin{cases} \lceil N/T \rceil & 1 \leq i < T \\ N - (T-1)\lceil N/T \rceil & i = T \end{cases}$$

We now derive the time complexity of our greedy algorithm (in section IV-B) with and without cost constraints. A preliminary step to Algorithm 2 involves

Algorithm 4 Node selection algorithm based on costs

Input: P
 $C_t \leftarrow$ CapEx available for time-step t
 $c_i \leftarrow$ cost of migrating node i
for $i \leftarrow 1$ to $|P|$ **do**
 $B(i) \leftarrow \frac{P_i}{c_i}$
end for
Sort P in descending order of corresponding entries in B
for $i \leftarrow 1$ to $|P|$ **do**
if $C_t \leq c_i$ **then**
 $C_t \leftarrow C_t - c_i$
Migrate node i
else
Break out of **for** loop
end if
end for

computation of all paths between every node pair, which can be achieved (say, using Floyd-Warshall Algorithm) in $O(N^3)$.

In Algorithm 2, the outer `for` loop is of the order of number of time-steps, i.e. $O(T)$, while the inner `for` loop runs across all unmigrated key-nodes, which is $O(N)$. The number of additional alternative paths made available by migration of a node, can be retrieved by scanning the list of precomputed paths between every node pair in constant time, given that we only consider the shortest hop-length paths between two nodes for the purpose of traffic engineering. The CapEx for a time-step can be computed based on the CapEx invested in the previous time-steps, resulting in $O(T)$.

In Algorithm 3, P can be sorted in $O(N \log N)$, whereas, the `for` loop runs over all nodes, i.e. $O(N)$. Thus, Algorithm 3 results in $O(N + N \log N)$.

In Algorithm 4, the first `for` loop runs over the list of unmigrated key-nodes, which is of the order of $O(N)$. The input element P is of length $O(N)$, and can be sorted in $O(N \log N)$. The second `for` loop also runs over the list of unmigrated key-nodes, resulting in $O(N)$ complexity. Thus, the order complexity of Algorithm 4 is $O(N + N \log N)$.

Cumulatively, the time complexity of Algorithm 2 can be approximated to $O(T \cdot N \log N)$, which comes in sharp contrast with the exponential complexity of the optimization algorithm.

V. SIMULATION AND RESULTS

For the performance evaluation, we developed a Java-based simulation framework. The topology used was the TA2 network (65 nodes, 108 links) from the SNDlib topology library [16]. We studied the migration profile of the nodes over 10 equal migration periods. We assigned uniformly random link weights assuming that the difference between any two link costs is below the inverse of the global maximum path length relative to the minimum link cost in the topology. This mimics OSPF's behavior in that between a node pair, no n -hop path should cost more than a $(n + 1)$ -hop path.

The migration sequence optimization was computed with the GUROBI optimizer on an Intel Core i7-3930K CPU (6 x 3.2 GHz) in less than 10 seconds, while the optimal traffic distribution computation took about 30 minutes (with an allowed MIP gap of 1%). Each of the plots in Figure 4 were averaged over 10 different random assignments of link weights and traffic profiles. Each traffic flow was uniformly distributed between 0 and 400 Mbit/s. We assumed links available in the following granularities (in Gbit/s): 1, 5, 10, 40, 100, 400, and 1000. To incorporate traffic growth during the course of the simulation, we set the growth factor of each flow in each time-step to a random value between 1.05 and 1.3, to ensure mean traffic growth of 20% in every time-step. We next demonstrate the importance




Fig. 3. Performance comparison of optimal, greedy, and random migration schedules without cost considerations

of effective migration sequence through two relevant metrics, namely, number of alternative available paths between a node pair, and network capacity savings.

Figure 3 plots the number of additional alternative paths resulting over time from various migration sequences - optimal, greedy, random. At the start of the simulation, none of the nodes had migrated, resulting in all curves starting at origin. Similarly, at the end of the simulation all nodes had migrated, resulting in a single end-point for all curves. The challenge thus remains in optimizing the *sequence* of migration, which we found to make a big difference in the number of alternative paths made available through SDN capabilities at migrated nodes, especially in the early stages. For the scenario without cost considerations, we restricted the number of nodes that can migrate in a time-step, namely, 7 nodes in each of the first 9 time-steps, 2 remaining nodes in the last time-step. For the case with cost constraints, we restricted the amount of CapEx investment per time-step. The cost of migrating a node was assumed proportional to its degree of connectivity in the network.

At the start of the simulation, we allot the network operator the total CapEx required for migration of all nodes in its network. It is remarkable to note how close the greedy algorithm comes to the optimal, both with and without cost constraints. The CapEx limit per time-step can also be seen from the step-like pattern of the curves with cost considerations, in contrast with the smooth curves without cost considerations. It is also noteworthy that the optimal (or greedy) sequence without cost constraints is higher than the corresponding curve with cost constraints. This is because the large number of alternative paths made available in the early stages results from a heavy CapEx investment in the early stages, which is checked with CapEx




Fig. 4. Network capacity savings of optimal, greedy, and random migration schedules

constraints. We also found similar behavior (not shown here for brevity), as in Figure 3, for Cost266 network (37 nodes, 57 links) as a medium-size topology and France network (25 nodes, 45 links) as a small-size topology from the SNDlib library [16].

Figure 4 plots the network capacity savings over time resulting from different migration sequences for the case with cost constraints. We note that the optimal migration sequence significantly outperforms the random migration sequence, while the greedy approach closely follows the optimal throughout. The drop observed in the curves at around the middle and end is due to the first-time use of a higher capacity link, 400 Gbit/s for the drop in the middle, and 1 Tbit/s for the drop at the end. These links are not fully utilized in the early stages.

In Figure 5, we compare the observed simulation run-times (in milliseconds, log-scale) of the migration sequence computation for the ILP and greedy approach for the case with cost constraints. Each data point in the plot was averaged over 10 runs. We observe that our heuristic increasingly outperforms the optimization model with increasing network size. This in turn, establishes the scalability of our algorithm, when compared to computation of the optimal migration sequence.

VI. CONCLUSIONS

In this paper, we presented optimization techniques and greedy algorithms to solve for the optimal migration sequence of SDN routers from the techno-economic perspective, and compared them analytically and empirically. The proposed approach included technological gains, as well as constraints on migration CapEx investment. In addition to the significantly




Fig. 5. Simulation runtime of Optimal and Greedy approach

lower time complexity, we observed that our greedy approach comes very close to the optimal, both in terms of number of alternative paths, as well as network capacity savings.

REFERENCES

- [1] ONF White paper. (2012, April 13) Software-defined networking: The new norm for networks. [Online]. Available: <https://www.opennetworking.org/images/stories/downloads/sdn-resources/white-papers/wp-sdn-newnorm.pdf>
- [2] M. Whitepaper. PCE - An Evolutionary Approach To SDN. [Online]. Available: <http://pcetosdn.com/downloads/PCE-Evolutionary-SDN-feb2012.pdf>
- [3] S. Sen, Y. Jin, R. Guerin, and K. Hosanagar, "Modeling the dynamics of network technology adoption and the role of converters," *Networking, IEEE/ACM Transactions on*, vol. 18, no. 6, pp. 1793–1805, 2010.
- [4] M. W. Macy and R. Willer, "From factors to actors: Computational sociology and agent-based modeling," *Annual Review of Sociology*, vol. 28, no. 1, pp. 143–166, 2002.
- [5] T. Trinh, L. Gyarmati, and G. Sallai, "Migrating to IPv6: A game-theoretic perspective," in *Local Computer Networks (LCN), 2010 IEEE 35th Conference on*, 2010, pp. 344–347.
- [6] P. Gill, M. Schapira, and S. Goldberg, "Let the market drive deployment: A strategy for transitioning to BGP security," in *Proc. of ACM SIGCOMM conference*. ACM, 2011, pp. 14–25.
- [7] S. Verbrugge, "Strategic planning of optical telecommunication networks in a dynamic and uncertain environment," Ph.D. dissertation, Ghent University, 2007.
- [8] S. Turk, R. Radeke, and R. Lehnert, "Network migration using ant colony optimization," in *Telecommunications Internet and Media Techno Economics (CTTE), IEEE Conference on*, 2010.
- [9] C. Kronberger, T. Schondienst, and D. A. Schupke, "Impact and handling of demand uncertainty in multiperiod planned networks," in *IEEE ICC*, 2011.
- [10] S. Rajagopalan, "Adoption timing of new equipment with another innovation anticipated," *Engineering Management, IEEE Transactions on*, vol. 46, no. 1, pp. 14–25, 1999.
- [11] B. Heller, R. Sherwood, and N. McKeown, "The controller placement problem," in *Proceedings of the first workshop on Hot topics in software defined networks*. ACM, 2012, pp. 7–12.
- [12] T. Das, M. Drogon, A. Jukan, and M. Hoffmann, "An agent-based modeling approach of network migration to new technologies," in *IEEE ICC*, 2013.

- [13] M. Caria, A. Jukan, and M. Hoffmann, "A performance study of network migration to SDN-enabled traffic engineering," in *IEEE Globecom*, December 2013.
- [14] Routeflow project, www.sites.google.com/site/routeflow/.
- [15] S. Das, G. Parulkar, and N. McKeown, "Why openflow/sdn can succeed where gmpls failed," in *European Conference and Exhibition on Optical Communication*. OSA, 2012.
- [16] SND library. [Online]. Available: <http://sndlib.zib.de>